

Air_Quality_Project

December 26, 2025

0.1 1. Introduction & Dataset Overview

- Source: Kaggle (city_day.csv – 29,531 rows, 16 columns)
- Goal: Clean messy real-world data, perform EDA, build ML model to predict AQI
- Problems: 15–61% missing values, wrong data types, outliers

0.2 2. Data Cleaning

0.2.1 Load and initial inspection

```
[194]: from google.colab import drive  
drive.mount('/content/drive')  
# Mount drive
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).

```
[195]: import pandas as pd
```

```
[196]: df = pd.read_csv("/content/drive/MyDrive/Air Quality India/city_day.csv")  
df.head()
```

```
[196]:      City        Date  PM2.5  PM10     NO    NO2   NOx  NH3     CO    SO2  \\\n0  Ahmedabad  2015-01-01    NaN    NaN  0.92  18.22  17.15  NaN  0.92  27.64\n1  Ahmedabad  2015-01-02    NaN    NaN  0.97  15.69  16.46  NaN  0.97  24.55\n2  Ahmedabad  2015-01-03    NaN    NaN  17.40  19.30  29.70  NaN  17.40  29.07\n3  Ahmedabad  2015-01-04    NaN    NaN  1.70  18.48  17.97  NaN  1.70  18.59\n4  Ahmedabad  2015-01-05    NaN    NaN  22.10  21.42  37.76  NaN  22.10  39.33\n\n          O3  Benzene  Toluene  Xylene  AQI  AQI_Bucket\n0  133.36     0.00     0.02     0.00  NaN        NaN\n1   34.06     3.68     5.50     3.77  NaN        NaN\n2   30.70     6.80    16.40     2.25  NaN        NaN\n3   36.08     4.43    10.14     1.00  NaN        NaN\n4   39.31     7.01    18.89     2.78  NaN        NaN
```

```
[197]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29531 entries, 0 to 29530
Data columns (total 16 columns):
 #   Column      Non-Null Count Dtype  
--- 
 0   City         29531 non-null   object  
 1   Date         29531 non-null   object  
 2   PM2.5        24933 non-null   float64 
 3   PM10         18391 non-null   float64 
 4   NO            25949 non-null   float64 
 5   NO2           25946 non-null   float64 
 6   NOx          25346 non-null   float64 
 7   NH3           19203 non-null   float64 
 8   CO            27472 non-null   float64 
 9   SO2           25677 non-null   float64 
 10  O3            25509 non-null   float64 
 11  Benzene       23908 non-null   float64 
 12  Toluene       21490 non-null   float64 
 13  Xylene        11422 non-null   float64 
 14  AQI           24850 non-null   float64 
 15  AQI_Bucket    24850 non-null   object  
dtypes: float64(13), object(3)
memory usage: 3.6+ MB

```

0.2.2 Fix data types (Date to datetime)

```
[198]: # Step: Convert the 'Date' column to the correct datetime type
df['Date'] = pd.to_datetime(df['Date'], errors='coerce')

# Let's check if it worked
print(df['Date'].dtype)
df['Date'].head(10)
```

datetime64[ns]

```
[198]: 0   2015-01-01
 1   2015-01-02
 2   2015-01-03
 3   2015-01-04
 4   2015-01-05
 5   2015-01-06
 6   2015-01-07
 7   2015-01-08
 8   2015-01-09
 9   2015-01-10
Name: Date, dtype: datetime64[ns]
```

```
[199]: # Let's check the current data types of all columns
print(df.dtypes)
```

```
City          object
Date         datetime64[ns]
PM2.5        float64
PM10         float64
NO           float64
NO2          float64
NOx          float64
NH3          float64
CO           float64
SO2          float64
O3            float64
Benzene      float64
Toluene      float64
Xylene       float64
AQI          float64
AQI_Bucket   object
dtype: object
```

0.2.3 Handle missing values (smart imputation)

- Numerical pollutants & AQI: median per city → global median fallback
- AQI_Bucket (categorical): most frequent value (mode)

```
[200]: # How many missing values in each column? (as a percentage - easier to understand)
missing_percent = (df.isnull().sum() / len(df)) * 100
missing_percent.round(2)
```

```
[200]: City      0.00
Date      0.00
PM2.5     15.57
PM10      37.72
NO        12.13
NO2       12.14
NOx       14.17
NH3        34.97
CO         6.97
SO2       13.05
O3        13.62
Benzene    19.04
Toluene    27.23
Xylene     61.32
AQI        15.85
AQI_Bucket 15.85
dtype: float64
```

```
[201]: # Step: Fill missing PM2.5 values with the median FOR EACH CITY SEPARATELY
# (straight from Pandas documentation - groupby + transform)

df['PM2.5'] = df.groupby('City')['PM2.5'].transform(lambda x: x.fillna(x.
    .median()))

# Let's check how many are still missing in PM2.5
print("Missing PM2.5 after filling per city:", df['PM2.5'].isnull().sum())
```

Missing PM2.5 after filling per city: 0

```
[202]: # Step: Fill missing PM10 values with the median FOR EACH CITY SEPARATELY
df['PM10'] = df.groupby('City')['PM10'].transform(lambda x: x.fillna(x.
    .median()))

# Let's check the result
print("Missing PM10 after filling per city:", df['PM10'].isnull().sum())
```

Missing PM10 after filling per city: 2009

```
[203]: # Step: Fill remaining missing PM10 values with the global median
global_median_pm10 = df['PM10'].median()

df['PM10'].fillna(global_median_pm10, inplace=True)

# Let's check the final result
print("Missing PM10 after second filling:", df['PM10'].isnull().sum())
```

Missing PM10 after second filling: 0

/tmp/ipython-input-99903386.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['PM10'].fillna(global_median_pm10, inplace=True)
```

```
[204]: # 1. First, median per city for NO2
df['NO2'] = df.groupby('City')['NO2'].transform(lambda x: x.fillna(x.median()))

# Check how many are left
```

```
print("Missing NO2 after median per city:", df['NO2'].isnull().sum())
```

Missing NO2 after median per city: 0

```
[205]: # 1. Median per city for NO
df['NO'] = df.groupby('City')['NO'].transform(lambda x: x.fillna(x.median()))

# Check how many are left
print("Missing NO after median per city:", df['NO'].isnull().sum())
```

Missing NO after median per city: 0

```
[206]: # 1. Median per city for NOx
df['NOx'] = df.groupby('City')['NOx'].transform(lambda x: x.fillna(x.median()))

# Check how many are left
print("Missing NOx after median per city:", df['NOx'].isnull().sum())
```

Missing NOx after median per city: 1169

```
[207]: # 2. Fill remaining missing NOx values with the global median
global_median_nox = df['NOx'].median()

df['NOx'].fillna(global_median_nox, inplace=True)

# Check the final result
print("Missing NOx after global median:", df['NOx'].isnull().sum())
```

Missing NOx after global median: 0

/tmp/ipython-input-1287824488.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['NOx'].fillna(global_median_nox, inplace=True)
```

```
[208]: # 1. Median per city for CO
df['CO'] = df.groupby('City')['CO'].transform(lambda x: x.fillna(x.median()))

# Check how many are left
```

```
print("Missing CO after median per city:", df['CO'].isnull().sum())
```

Missing CO after median per city: 0

```
[209]: # 1. Median per city for S02  
df['S02'] = df.groupby('City')['S02'].transform(lambda x: x.fillna(x.median()))  
  
# Check how many are left  
print("Missing S02 after median per city:", df['S02'].isnull().sum())
```

Missing S02 after median per city: 0

```
[210]: # 1. Median per city for O3  
df['O3'] = df.groupby('City')['O3'].transform(lambda x: x.fillna(x.median()))  
  
# Check how many are left  
print("Missing O3 after median per city:", df['O3'].isnull().sum())
```

Missing O3 after median per city: 162

```
[211]: # 2. Fill remaining missing O3 values with the global median  
global_median_o3 = df['O3'].median()  
  
df['O3'].fillna(global_median_o3, inplace=True)  
  
# Check the final result  
print("Missing O3 after global median:", df['O3'].isnull().sum())
```

Missing O3 after global median: 0

/tmp/ipython-input-904117558.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['O3'].fillna(global_median_o3, inplace=True)
```

```
[212]: # 1. Median per city for NH3  
df['NH3'] = df.groupby('City')['NH3'].transform(lambda x: x.fillna(x.median()))  
  
# Check how many are left
```

```
print("Missing NH3 after median per city:", df['NH3'].isnull().sum())
```

Missing NH3 after median per city: 2009

```
[213]: # 2. Fill remaining missing NH3 values with the global median
global_median_nh3 = df['NH3'].median()

df['NH3'].fillna(global_median_nh3, inplace=True)

# Check the final result
print("Missing NH3 after global median:", df['NH3'].isnull().sum())
```

Missing NH3 after global median: 0

/tmp/ipython-input-2115361960.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['NH3'].fillna(global_median_nh3, inplace=True)
```

```
[214]: # 1. Median per city for Benzene
df['Benzene'] = df.groupby('City')['Benzene'].transform(lambda x: x.fillna(x.median()))

# Check how many are left
print("Missing Benzene after median per city:", df['Benzene'].isnull().sum())
```

Missing Benzene after median per city: 2732

```
[215]: # 2. Fill remaining missing Benzene values with the global median
global_median_benzene = df['Benzene'].median()

df['Benzene'].fillna(global_median_benzene, inplace=True)

# Check the final result
print("Missing Benzene after global median:", df['Benzene'].isnull().sum())
```

Missing Benzene after global median: 0

/tmp/ipython-input-1495005828.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace

method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Benzene'].fillna(global_median_benzene, inplace=True)
```

```
[216]: # 1. Median per city for Toluene
df['Toluene'] = df.groupby('City')['Toluene'].transform(lambda x: x.fillna(x.median()))

# Check how many are left
print("Missing Toluene after median per city:", df['Toluene'].isnull().sum())
```

Missing Toluene after median per city: 4010

```
[217]: # 2. Fill remaining missing Toluene values with the global median
global_median_toluene = df['Toluene'].median()

df['Toluene'].fillna(global_median_toluene, inplace=True)

# Check the final result
print("Missing Toluene after global median:", df['Toluene'].isnull().sum())
```

Missing Toluene after global median: 0

/tmp/ipython-input-679443913.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Toluene'].fillna(global_median_toluene, inplace=True)
```

```
[218]: # 1. Median per city for Xylene
df['Xylene'] = df.groupby('City')['Xylene'].transform(lambda x: x.fillna(x.median()))
```

```
# Check how many are left
print("Missing Xylene after median per city:", df['Xylene'].isnull().sum())
```

Missing Xylene after median per city: 13047

```
[219]: # 2. Fill remaining missing Xylene values with the global median
global_median_xylene = df['Xylene'].median()

df['Xylene'].fillna(global_median_xylene, inplace=True)

# Check the final result
print("Missing Xylene after global median:", df['Xylene'].isnull().sum())
```

Missing Xylene after global median: 0

/tmp/ipython-input-2455260267.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['Xylene'].fillna(global_median_xylene, inplace=True)
```

```
[220]: # 1. Median per city for AQI
df['AQI'] = df.groupby('City')['AQI'].transform(lambda x: x.fillna(x.median()))

# Check how many are left
print("Missing AQI after median per city:", df['AQI'].isnull().sum())
```

Missing AQI after median per city: 0

```
[221]: # Fill missing AQI_Bucket values with the most frequent category
most_common_bucket = df['AQI_Bucket'].mode()[0] # mode() returns a series, we\u202d take the first element

df['AQI_Bucket'].fillna(most_common_bucket, inplace=True)

# Check the final result
print("Missing AQI_Bucket after filling:", df['AQI_Bucket'].isnull().sum())

# Bonus: see the category distribution
```

```
print("\nAQI_Bucket category distribution:")
print(df['AQI_Bucket'].value_counts())
```

Missing AQI_Bucket after filling: 0

AQI_Bucket category distribution:

```
AQI_Bucket
Moderate      13510
Satisfactory   8224
Poor          2781
Very Poor     2337
Good           1341
Severe         1338
Name: count, dtype: int64
```

/tmp/ipython-input-2445940677.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['AQI_Bucket'].fillna(most_common_bucket, inplace=True)
```

[222]: # How many missing values in each column?
df.isnull().sum()

[222]:

City	0
Date	0
PM2.5	0
PM10	0
NO	0
NO2	0
NOx	0
NH3	0
CO	0
S02	0
O3	0
Benzene	0
Toluene	0
Xylene	0
AQI	0
AQI_Bucket	0

```
dtype: int64
```

```
[223]: # Save the cleaned dataset to a file on Google Drive
from google.colab import drive
drive.mount('/content/drive')

# Save the file (change path if desired)
df.to_csv('/content/drive/MyDrive/Air Quality India/air_quality_india_cleaned.
         ↪csv', index=False)

print("Cleaned file saved to your Google Drive!")
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).

Cleaned file saved to your Google Drive!

0.3 3. Exploratory Data Analysis (EDA)


0.3.1 3.1 PM2.5 trend in Delhi (2015–2020)

Shows massive winter spikes due to stubble burning, heating and inversion

```
[224]: import matplotlib.pyplot as plt

# Filter for Delhi only and sort by date
delhi_df = df[df['City'] == 'Delhi'].sort_values('Date')

# Line plot of PM2.5 over time
plt.figure(figsize=(14, 6))
plt.plot(delhi_df['Date'], delhi_df['PM2.5'], color='red', linewidth=1)
plt.title('PM2.5 in Delhi (2015-2020) - after data cleaning', fontsize=16)
plt.xlabel('Date')
plt.ylabel('PM2.5 (µg/m³)')
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```



0.3.2 3.2 AQI distribution across major cities

Delhi and northern cities have significantly worse air quality

```
[225]: import seaborn as sns
import matplotlib.pyplot as plt

# Average AQI for major cities
major_cities = ['Delhi', 'Mumbai', 'Bengaluru', 'Kolkata', 'Chennai', 'Hyderabad', 'Ahmedabad', 'Lucknow']


city_aqi = df[df['City'].isin(major_cities)].groupby('City')['AQI'].mean().
    sort_values(ascending=False)

# Boxplot (shows distribution better than just the average)
plt.figure(figsize=(12, 6))
sns.boxplot(x='City', y='AQI', data=df[df['City'].isin(major_cities)], order=city_aqi.index, palette='Reds_r')
plt.title('AQI Distribution in Major Indian Cities (2015-2020)', fontsize=16)
plt.xlabel('City')
plt.ylabel('AQI')
plt.xticks(rotation=45)
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```

/tmp/ipython-input-288516690.py:11: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='City', y='AQI', data=df[df['City'].isin(major_cities)],
order=city_aqi.index, palette='Reds_r')
```



0.3.3 3.3 Correlation heatmap


PM2.5, PM10 and CO are the strongest drivers of AQI

```
[226]: import seaborn as sns
import matplotlib.pyplot as plt

# Select only numerical columns (pollutants + AQI)
pollutants_cols = ['PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3', 'CO', 'SO2', 'O3', 'Benzene', 'Toluene', 'Xylene', 'AQI']

corr_matrix = df[pollutants_cols].corr()

# Heatmap
plt.figure(figsize=(12, 10))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.5, center=0)
plt.title('Correlation between Pollutants and AQI (2015-2020)', fontsize=16)
plt.tight_layout()
plt.show()
```



0.4 4. Machine Learning – Predicting AQI

0.4.1 Model: Random Forest Regressor

- Features: 12 pollutants
- Train/test split: 80/20
- Result: $R^2 = 0.9$ on test set (excellent performance)

```
[227]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import r2_score, mean_absolute_error
import numpy as np

# Data preparation - using pollutants to predict AQI
features = ['PM2.5', 'PM10', 'NO', 'NO2', 'NOx', 'NH3', 'CO', 'SO2', 'O3', 'Benzene', 'Toluene', 'Xylene']
```

```

X = df[features]
y = df['AQI']

# Split into train/test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Random Forest Model
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

# Prediction and results
y_pred = model.predict(X_test)
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)

print(f"R2 score: {r2:.4f} ← closer to 1.0 is better (1.0 = perfect prediction)")
print(f"Mean Absolute Error: {mae:.2f} ← average error in AQI units")

# Feature importance - which pollutants are most important for the model
importances = model.feature_importances_
feat_imp = sorted(zip(features, importances), key=lambda x: x[1], reverse=True)
print("\nMost important pollutants according to the model:")
for feat, imp in feat_imp:
    print(f"{feat}: {imp:.4f}")

```

R² score: 0.9016 ← closer to 1.0 is better (1.0 = perfect prediction)
 Mean Absolute Error: 19.82 ← average error in AQI units

Most important pollutants according to the model:
 PM2.5: 0.4277
 CO: 0.4075
 NO: 0.0469
 PM10: 0.0389
 SO2: 0.0143
 O₃: 0.0139
 NOx: 0.0132
 NO2: 0.0114
 Toluene: 0.0091
 Xylene: 0.0063
 Benzene: 0.0061
 NH3: 0.0047